

QAA – DATA TRANSFORMATION

QA

USING PYTHON



- Using python


USING PYTHON

- Data Transformation

Importing packages

Import pandas and numpy

```
► import pandas as pd  
    import numpy as np
```



Loading data

Hard-coded data entry

```
N data = {  
    'Name': ['Alice', 'Bob Gunner', 'Charlie', np.nan, 'Eve', 'Frank', 'Grace', 'Alice'],  
    'Age': [25, np.nan, 35, 40, 50, 'Unknown', 30, 25],  
    'Gender': ['Female', 'Male', np.nan, 'Female', 'Other', 'Male', 'Female', 'Female'],  
    'Salary': [50000, 52000, 55000, np.nan, 48000, 'N/A', 51000, 50000],  
    'Email': ['alice@email.com', 'bob@email', 'charlie--email.com', '', 'eve@email.com', '', 'grace@email.com', 'alice@email.  
    'Retired':[False,True,False,False,False,False,False]  
}  
  
N df = pd.DataFrame(data)  
print(df)
```

	Name	Age	Gender	Salary	Email	Retired
0	Alice	25	Female	50000	alice@email.com	False
1	Bob Gunner	NaN	Male	52000	bob@email	True
2	Charlie	35	NaN	55000	charlie--email.com	False
3	NaN	40	Female	NaN		False
4	Eve	50	Other	48000	eve@email.com	False
5	Frank	Unknown	Male	N/A		False
6	Grace	30	Female	51000	grace@email.com	False
7	Alice	25	Female	50000	alice@email.com	False

Loading data

CSV Loading

```
| df_csv = pd.read_csv("Seattle_cycles_station.csv",sep=',')
```

```
| df_csv.head(5)
```


	station_id	name	lat	long	install_date	install_dockcount	modification_date	current_dockcount	decommission_date
0	BT-01	3rd Ave & Broad St	47.618418	-122.350964	10/13/2014	18	NaN	18	NaN
1	BT-03	2nd Ave & Vine St	47.615829	-122.348564	10/13/2014	16	NaN	16	NaN
2	BT-04	6th Ave & Blanchard St	47.616094	-122.341102	10/13/2014	16	NaN	16	NaN
3	BT-05	2nd Ave & Blanchard St	47.613110	-122.344208	10/13/2014	14	NaN	14	NaN
4	CBD-03	7th Ave & Union St	47.610731	-122.332447	10/13/2014	20	NaN	20	NaN

EXERCISE 1: LOADING THE DATAFRAME



- Install pandas into a notebook
- Load the data from a csv file called seattle_cycles_station.csv into a dataframe called df_cycle_station
- View the dataframe using the print function


Using tail(n)



df_csv.tail(8)

	station_id	name	lat	long	install_date	install_dockcount	modification_date	current_dockcount	decommission_date
50	UW-07	UW Intramural Activities Building	47.653713	-122.302162	10/13/2014	20	2/20/2015	14	NaN
51	UW-10	UW Magnuson Health Sciences Center Rotunda / C...	47.650725	-122.311188	10/13/2014	16	NaN	16	NaN
52	WF-01	Pier 69 / Alaskan Way & Clay St	47.614315	-122.354093	10/13/2014	18	NaN	24	NaN
53	WF-04	Seattle Aquarium / Alaskan Way S & Elliott Bay...	47.607702	-122.341650	10/13/2014	18	NaN	18	NaN
54	CH-16	Broadway and E Denny Way	47.618640	-122.320777	3/18/2016	18	NaN	18	NaN
55	SLU-22	Thomas St & 5th Ave N	47.620879	-122.347377	7/3/2016	18	NaN	18	NaN
56	UW-11	NE Pacific St/UW Medical Center	47.649952	-122.306263	10/29/2015	16	NaN	16	NaN
57	WF-03	Pier 66 / Alaskan Way & Bell St	47.611370	-122.348702	8/9/2016	18	NaN	18	NaN

Using head(n)



df_csv.tail(8)


	station_id	name	lat	long	install_date	install_dockcount	modification_date	current_dockcount	decommission_date
50	UW-07	UW Intramural Activities Building	47.653713	-122.302162	10/13/2014	20	2/20/2015	14	NaN
51	UW-10	UW Magnuson Health Sciences Center Rotunda / C...	47.650725	-122.311188	10/13/2014	16	NaN	16	NaN
52	WF-01	Pier 69 / Alaskan Way & Clay St	47.614315	-122.354093	10/13/2014	18	NaN	24	NaN
53	WF-04	Seattle Aquarium / Alaskan Way S & Elliott Bay...	47.607702	-122.341650	10/13/2014	18	NaN	18	NaN
54	CH-16	Broadway and E Denny Way	47.618640	-122.320777	3/18/2016	18	NaN	18	NaN
55	SLU-22	Thomas St & 5th Ave N	47.620879	-122.347377	7/3/2016	18	NaN	18	NaN
56	UW-11	NE Pacific St/UW Medical Center	47.649952	-122.306263	10/29/2015	16	NaN	16	NaN
57	WF-03	Pier 66 / Alaskan Way & Bell St	47.611370	-122.348702	8/9/2016	18	NaN	18	NaN

EXERCISE 2: REVIEWING THE FIRST AND LAST ROWS



- Showing the LAST and HEAD functions to review the content of the dataframe (5 rows only)

Showing the columns and data types



```
▶ print(df_csv.columns)
```

```
Index(['station_id', 'name', 'lat', 'long', 'install_date',
       'install_dockcount', 'modification_date', 'current_dockcount',
       'decommission_date'],
      dtype='object')
```

```
▶ print(df_csv.dtypes)
```


station_id	object
name	object
lat	float64
long	float64
install_date	object
install_dockcount	int64
modification_date	object
current_dockcount	int64
decommission_date	object
dtype:	object

EXERCISE 3: REVIEWING THE DATA TYPES



- Write the code to review the columns available in the df_cycle_station dataframe
- Write the code to review the data types of the dataframe

Showing number of cells, rows and columns



How many cells does the dataframe have?

```
▶ print(df_cycle_station.size)
```

580

How the number of rows and coloumns?

```
▶ df_cycle_station.shape[0]
```

58

```
▶ df_cycle_station.shape[1]
```


10

EXERCISE 4: REVIEWING THE DATAFRAME SIZE



- Write the code to review the columns available in the df_cycle_station dataframe
- Write the code to review the data types of the dataframe

Replacing certain values



Replacing the N/A values with a set value

```
df_filled = df.fillna({'Name': 'Unknown', 'Gender': 'Unknown'})  
print(df_filled)
```

	Name	Age	Gender	Salary	Email
0	Alice	25	Female	50000	alice@email.com
1	Bob Gunner	NaN	Male	52000	bob@email
2	Charlie	35	Unknown	55000	charlie--email.com
3	Unknown	40	Female	NaN	
4	Eve	50	Other	48000	eve@email.com
5	Frank	Unknown	Male	N/A	
6	Grace	30	Female	51000	grace@email.com
7	Alice	25	Female	50000	alice@email.com

replace NaN with set values


```
df.replace({'Age': 'Unknown', 'Salary': 'N/A'}, np.nan, inplace=True)  
print(df)
```

	Name	Age	Gender	Salary	Email
0	Alice	25.0	Female	50000.0	alice@email.com
1	Bob Gunner	NaN	Male	52000.0	bob@email
2	Charlie	35.0	NaN	55000.0	charlie--email.com
3	NaN	40.0	Female	NaN	
4	Eve	50.0	Other	48000.0	eve@email.com
5	Frank	NaN	Male	NaN	
6	Grace	30.0	Female	51000.0	grace@email.com
7	Alice	25.0	Female	50000.0	alice@email.com

Calculate a replacement value

Replacing age and salary with a calculated replacement

```
▶ age_median = df['Age'].median(skipna=True)
   salary_mean = df['Salary'].mean(skipna=True)
   df['Age'].fillna(age_median, inplace=True)
   df['Salary'].fillna(salary_mean, inplace=True)
   print(df)
```



	Name	Age	Gender	Salary	Email
0	Alice	25.0	Female	50000.0	alice@email.com
1	Bob Gunner	32.5	Male	52000.0	bob@email
2	Charlie	35.0	NaN	55000.0	charlie--email.com
3	NaN	40.0	Female	51000.0	
4	Eve	50.0	Other	48000.0	eve@email.com
5	Frank	32.5	Male	51000.0	
6	Grace	30.0	Female	51000.0	grace@email.com
7	Alice	25.0	Female	50000.0	alice@email.com

EXERCISE 5: CLEANING DATA



- Write the code to review the decommission_date within the dataframe
- Write the code to replace any blank values with '12/31/2030'
- Review the column again to check the values have been changed
- Write the code to replace all decommission_date rows that have '12/31/2030' with np.nan

Changing data types

Change the data types of the numeric fields

```
df['Age'] = df['Age'].astype(int)
df['Salary'] = df['Salary'].astype(float)
print(df)
```

	Name	Age	Gender	Salary	Email
0	Alice	25	Female	50000.0	alice@email.com
1	Bob Gunner	32	Male	52000.0	bob@email
2	Charlie	35	NaN	55000.0	charlie--email.com
3	NaN	40	Female	51000.0	
4	Eve	50	Other	48000.0	eve@email.com
5	Frank	32	Male	51000.0	
6	Grace	30	Female	51000.0	grace@email.com
7	Alice	25	Female	50000.0	alice@email.com

Working with Dates

Working with dates

```
▶ date_example_df = pd.DataFrame({  
    'Name': ['Alice', 'Bob Newhart'],  
    'DOB_String': ['2021-08-10 15:30:45', '2021-08-10 15:30:45']  
})  
date_example_df['DOB_DateTime'] = pd.to_datetime(date_example_df['DOB_String'])  
  
print(date_example_df)
```

	Name	DOB_String	DOB_DateTime
0	Alice	2021-08-10 15:30:45	2021-08-10 15:30:45
1	Bob Newhart	2021-08-10 15:30:45	2021-08-10 15:30:45

```
▶ # Example: '2021-08-10 15:30:45'  
date_example_df['Year'] = date_example_df['DOB_DateTime'].dt.year  
date_example_df['Month'] = date_example_df['DOB_DateTime'].dt.month  
date_example_df['day'] = date_example_df['DOB_DateTime'].dt.day  
print(date_example_df)
```

	Name	DOB_String	DOB_DateTime	Year	Month	day
0	Alice	2021-08-10 15:30:45	2021-08-10 15:30:45	2021	8	10
1	Bob Newhart	2021-08-10 15:30:45	2021-08-10 15:30:45	2021	8	10

```
▶ date_example_df['hour'] = date_example_df['DOB_DateTime'].dt.hour  
date_example_df['minute'] = date_example_df['DOB_DateTime'].dt.minute  
date_example_df['second'] = date_example_df['DOB_DateTime'].dt.second  
print(date_example_df)
```

	Name	DOB_String	DOB_DateTime	Year	Month	day	hour	minute	second
0	Alice	2021-08-10 15:30:45	2021-08-10 15:30:45	2021	8	10	15	30	45
1	Bob Newhart	2021-08-10 15:30:45	2021-08-10 15:30:45	2021	8	10	15	30	45

EXERCISE 6: CHANGING DATA TYPES



- Write the code to review the data types of all columns in the dataframe
- Change the data type of the following columns:
 - Station_id and name to string
 - Lat and long to float64
- Write the code to review the data types of the changed columns in the dataframe
- Use `pd.to_datetime(df['column'])` to change the install_date, modification_date and decomission_date to a date column

Normalize data

Normalise the numeric columns


```
df['Salary_Normalized'] = (df['Salary'] - df['Salary'].min()) / (df['Salary'].max() - df['Salary'].min())
df['Age_Normalized'] = (df['Age'] - df['Age'].min()) / (df['Age'].max() - df['Age'].min())
print(df)
```

	Name	Age	Gender	Salary	Email	Salary_Normalized
0	Alice	25	Female	50000.0	alice@email.com	0.285714
1	Bob Gunner	32	Male	52000.0	bob@email	0.571429
2	Charlie	35	Nan	55000.0	charlie--email.com	1.000000
3	NaN	40	Female	51000.0		0.428571
4	Eve	50	Other	48000.0	eve@email.com	0.000000
5	Frank	32	Male	51000.0		0.428571
6	Grace	30	Female	51000.0	grace@email.com	0.428571
7	Alice	25	Female	50000.0	alice@email.com	0.285714

Age_Normalized

0	0.00
1	0.28
2	0.40
3	0.60
4	1.00
5	0.28
6	0.20
7	0.00

Placing in categories (Binning)




Binning of numeric values to group

```
bins = [0, 30, 40, 50, 60]
labels = ['20s', '30s', '40s', '50s']
df['Age Group'] = pd.cut(df['Age'], bins=bins, labels=labels, right=False)
print(df)
```

	Name	Age	Gender	Salary	Email	Salary_Normalized
0	Alice	25	Female	50000.0	alice@email.com	0.285714
1	Bob Gunner	32	Male	52000.0	bob@email	0.571429
2	Charlie	35	NaN	55000.0	charlie--email.com	1.000000
3	NaN	40	Female	51000.0		0.428571
4	Eve	50	Other	48000.0	eve@email.com	0.000000
5	Frank	32	Male	51000.0		0.428571
6	Grace	30	Female	51000.0	grace@email.com	0.428571
7	Alice	25	Female	50000.0	alice@email.com	0.285714

	Age_Normalized	Age Group
0	0.00	20s
1	0.28	30s
2	0.40	30s
3	0.60	40s
4	1.00	50s
5	0.28	30s
6	0.20	30s
7	0.00	20s

Create category columns



Categorise the gender column (note the extra columns with boolean results)

```
# gender_encoded = pd.get_dummies(df['Gender'], prefix='Gender')
df = pd.concat([df, gender_encoded], axis=1)
print(df)
```


	Name	Age	Gender	Salary	Email	Salary_Normalized	\
0	Alice	25	Female	50000.0	alice@email.com	0.285714	
1	Bob Gunner	32	Male	52000.0	bob@email	0.571429	
2	Charlie	35	NaN	55000.0	charlie--email.com	1.000000	
3	Nan	40	Female	51000.0		0.428571	
4	Eve	50	Other	48000.0	eve@email.com	0.000000	
5	Frank	32	Male	51000.0		0.428571	
6	Grace	30	Female	51000.0	grace@email.com	0.428571	
7	Alice	25	Female	50000.0	alice@email.com	0.285714	

	Age_Normalized	Age	Group	Gender_Female	Gender_Male	Gender_Other
0	0.00	20s		1	0	0
1	0.28	30s		0	1	0
2	0.40	30s		0	0	0
3	0.60	40s		1	0	0
4	1.00	50s		0	0	1
5	0.28	30s		0	1	0
6	0.20	30s		1	0	0
7	0.00	20s		1	0	0

EXERCISE 7: NORMALISING G, BINNING AND CATEGORIES FROM VALUES

- Normalise the lat and long columns into the columns lat_normal and long_normal
- List the unique values from the install_dockcount and current_dockcount columns
- Create bins for the dockcount as follows:
 - 0,1,10,15,24,30,100
 - Empty, Very-Small, Small, Medium, Large, Extra-Large
- Bin the install_dockcount column of the dataframe into the column install_group
- Bin the current_dockcount column of the dataframe into the column current_group
- Display the dataframe columns (install_dockcount, install_group, current_dockcount, current_group)

Function for checking email has @




create a function to validate email address (@) and replace if needed

```
# extended reading on functions
def is_valid_email(email):
    if '@' in email:
        return email
    else:
        return np.nan

df['Email'] = df['Email'].apply(is_valid_email)
print(df['Email'])
```

```
0    alice@email.com
1          bob@email
2              NaN
3              NaN
4    eve@email.com
5              NaN
6  grace@email.com
7  alice@email.com
Name: Email, dtype: object
```

Working with strings



Splitting the address

```
▶ df_cycle_station['left'] = df_cycle_station['name'].str.split('&').str[0]
   print(df_cycle_station)
```

Testing strings

```
▶ df_cycle_station['Length_Name'] = df_cycle_station['name'].str.len()
   df_cycle_station['name'] = df_cycle_station['name'].str.strip()
   df_cycle_station['Length_NameAfterStrip'] = df_cycle_station['name'].str.len()
   df_cycle_station
```

Checking lengths strings within columns

Show the lengths of two text fields (before and after strip)

```
df['Length_Email'] = df['Email'].str.len()  
df['Length_Name'] = df['Name'].str.len()  
print(df)
```

	Name	Age	Gender	Salary	Email	Salary_Normalized
0	Alice	25	Female	50000.0	alice@email.com	0.285714
1	Bob Gunner	32	Male	52000.0	bob@email	0.571429
2	Charlie	35	NAN	55000.0	charlie--email.com	1.000000
3	NaN	40	Female	51000.0		0.428571
4	Eve	50	Other	48000.0	eve@email.com	0.000000
5	Frank	32	Male	51000.0		0.428571
6	Grace	30	Female	51000.0	grace@email.com	0.428571
7	Alice	25	Female	50000.0	alice@email.com	0.285714

	Age_Normalized	Age	Group	Gender_Female	Gender_Male	Gender_Other
0	0.00		20s	1	0	0
1	0.28		30s	0	1	0
2	0.40		30s	0	0	0
3	0.60		40s	1	0	0
4	1.00		50s	0	0	1
5	0.28		30s	0	1	0
6	0.20		30s	1	0	0
7	0.00		20s	1	0	0

	Length_Email	Length_Name
0	15	5.0
1	9	10.0
2	18	12.0
3	0	NaN
4	13	13.0
5	0	5.0
6	15	5.0
7	15	5.0

Replace cleansed column

```
df['Name'] = df['Name'].str.strip()  
df['Length_NameAfterStrip'] = df['Name'].str.len()  
print(df)
```


	Name	Age	Gender	Salary	Email	Salary_Normalized	\
0	Alice	25	Female	50000.0	alice@email.com	0.285714	
1	Bob Gunner	32	Male	52000.0	bob@email	0.571429	
2	Charlie	35	NaN	55000.0	charlie--email.com	1.000000	
3	NaN	40	Female	51000.0		0.428571	
4	Eve	50	Other	48000.0	eve@email.com	0.000000	
5	Frank	32	Male	51000.0		0.428571	
6	Grace	30	Female	51000.0	grace@email.com	0.428571	
7	Alice	25	Female	50000.0	alice@email.com	0.285714	

	Age_Normalized	Age	Group	Gender_Female	Gender_Male	Gender_Other	\
0	0.00		20s	1	0	0	
1	0.28		30s	0	1	0	
2	0.40		30s	0	0	0	
3	0.60		40s	1	0	0	
4	1.00		50s	0	0	1	
5	0.28		30s	0	1	0	
6	0.20		30s	1	0	0	
7	0.00		20s	1	0	0	

	Length_Email	Length_Name	Length_NameAfterStrip	
0	15	5.0	5.0	
1	9	10.0	10.0	
2	18	12.0	7.0	
3	0	NaN	NaN	
4	13	13.0	3.0	
5	0	5.0	5.0	
6	15	5.0	5.0	
7	15	5.0	5.0	



Standardising strings



Using the title, upper and lower function to format the strings

```
df['NameTitle'] = df['Name'].str.title()
print(df['NameTitle'])
```

```
0      Alice
1    Bob Gunner
2     Charlie
3       NaN
4      Eve
5     Frank
6    Grace
7      Alice
Name: NameTitle, dtype: object
```

```
df['NameUpper'] = df['Name'].str.upper()
print(df['NameUpper'])
```

```
0      ALICE
1    BOB GUNNER
2     CHARLIE
3       NaN
4      EVE
5     FRANK
6    GRACE
7      ALICE
Name: NameUpper, dtype: object
```

```
df['NameLower'] = df['Name'].str.lower()
print(df['NameLower'])
```

```
0      alice
1    bob gunner
2     charlie
3       NaN
4      eve
5     frank
6    grace
7      alice
Name: NameLower, dtype: object
```

EXERCISE 8: UPDATING STRINGS



- Split the name column at the character & placing the characters to the left into the column shorten_name
- Split the station_id column at the character - placing the characters to the left into the column code
- Review the two columns (shorten_name, code0)
- Write code to:
 - Check the length of the name column (into length_name)
 - Clean the name column of any additional trailing spaces using strip (replacing the current name)
 - Check the length of the name column (into length_afterstrip)
 - Review the name, length_name, length_afterstrip

Dropping a column

```
df.drop(["NameUpper","NameTitle"], axis = 1,inplace = True)  
df.reset_index(drop=True)
```

```
print(df)
```

	Name	Age	Gender	Salary	Email	Retired	\
0	Alice	25	Female	50000.0	alice@email.com	False	
1	Bob Gunner	32	Male	52000.0	bob@email	True	
2	Charlie	35	Nan	55000.0	NaN	False	
3	NaN	40	Female	51000.0	NaN	False	
4	Eve	50	Other	48000.0	eve@email.com	False	
5	Frank	32	Male	51000.0	NaN	False	
6	Grace	30	Female	51000.0	grace@email.com	False	

	Salary_Normalized	Age_Normalized	Age	Group	Gender_Female	Gender_Male	\
0	0.285714	0.00	20s		1	0	
1	0.571429	0.28	30s		0	1	
2	1.000000	0.40	30s		0	0	
3	0.428571	0.60	40s		1	0	
4	0.000000	1.00	50s		0	0	
5	0.428571	0.28	30s		0	1	
6	0.428571	0.20	30s		1	0	

	Gender_Other	Length_Email	Length_Name	Length_NameAfterStrip	NameLower	
0	0	15	5.0	5.0	alice	
1	0	9	10.0	10.0	bob gunner	
2	0	18	12.0	7.0	charlie	
3	0	0	NaN	NaN	NaN	
4	1	13	13.0	3.0	eve	
5	0	0	5.0	5.0	frank	
6	0	15	5.0	5.0	grace	

EXERCISE 9:

REMOVE COLUMNS NOT REQUIRED FROM DATAFRAME



- Review the shape of dimension 1 of the dataframe
- Use the drop function to remove the title, lower, upper columns individually
- Review the shape of dimension 1 of the dataframe (and see the difference)
- Use the drop function to remove the length_name and length_after strip columns (as a single line of code)


Summarising data

Summarisation

```
▶ print(df.groupby('Gender').agg({'Age': ['mean', 'median','std','var','sum','first','last','min','max']}))  
print(df.groupby('Gender').agg({'Age': ['count', 'mean', 'nunique','unique']}))
```

Gender	Age									
	mean	median	std	var	sum	first	last	min	max	
Female	31.666667	30.0	7.637626	58.333333	95	25	30	25	40	
Male	32.000000	32.0	0.000000	0.000000	64	32	32	32	32	
Other	50.000000	50.0	NaN	NaN	50	50	50	50	50	
Gender	Age									
	count	mean	nunique	unique						
Female	3	31.666667	3	[25, 40, 30]						
Male	2	32.000000	1	[32]						
Other	1	50.000000	1	[50]						

Using describe



df.describe()


	Age	Salary	Salary_Normalized	Age_Normalized	Gender_Female	Gender_Male	Gender_Other	Length_Email	Length_Name	Length_Na
count	7.000000	7.000000	7.000000	7.000000	7.000000	7.000000	7.000000	7.000000	7.000000	6.000000
mean	34.857143	51142.857143	0.448980	0.394286	0.428571	0.285714	0.142857	10.000000	8.333333	
std	8.091736	2115.700942	0.302243	0.323669	0.534522	0.487950	0.377964	7.348469	3.777124	
min	25.000000	48000.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	5.000000	
25%	31.000000	50500.000000	0.357143	0.240000	0.000000	0.000000	0.000000	4.500000	5.000000	
50%	32.000000	51000.000000	0.428571	0.280000	0.000000	0.000000	0.000000	13.000000	7.500000	
75%	37.500000	51500.000000	0.500000	0.500000	1.000000	0.500000	0.000000	15.000000	11.500000	
max	50.000000	55000.000000	1.000000	1.000000	1.000000	1.000000	1.000000	18.000000	13.000000	

EXERCISE 10: SUMMARISE AND DESCRIBE



- Retrieve the following calculations
- (mean','median','std','var','sum','first','last','min','max','count', 'mean', 'nunique','unique
-)
- of the current_dockcount by code
- Use the describe function to review calculations for all numeric columns

Melting (unpivoting) Data



Moving columns to rows (unpivot in PBI)

```
df_melted = pd.melt(df, id_vars=['Name'], value_vars=['Age', 'Salary'])
print(df_melted)
```


	Name	variable	value
0	Alice	Age	25.0
1	Bob Gunner	Age	32.0
2	Charlie	Age	35.0
3	NaN	Age	40.0
4	Eve	Age	50.0
5	Frank	Age	32.0
6	Grace	Age	30.0
7	Alice	Salary	50000.0
8	Bob Gunner	Salary	52000.0
9	Charlie	Salary	55000.0
10	NaN	Salary	51000.0
11	Eve	Salary	48000.0
12	Frank	Salary	51000.0
13	Grace	Salary	51000.0

Pivoting data

Moving rows to columns (pivot in PBI)


```
df_pivoted = df.pivot(index='Name', columns='Gender', values='Salary')
print(df_pivoted)
```

Gender	NaN	Female	Male	Other
Name				
NaN	NaN	51000.0	NaN	NaN
Alice	NaN	50000.0	NaN	NaN
Bob Gunner	NaN	NaN	52000.0	NaN
Charlie	55000.0	NaN	NaN	NaN
Eve	NaN	NaN	NaN	48000.0
Frank	NaN	NaN	51000.0	NaN
Grace	NaN	51000.0	NaN	NaN





Concat two data frames



		Name	Age	Gender	Salary	Email	DataGroup
0		Alice	25	Female	50000	alice@email.com	Company_A
1	Bob	Newhart	NaN	Male	52000	bob@email	Company_A
2		Charlie	35	NaN	55000	charlie--email.com	Company_A
3		NaN	40	Female	NaN		Company_A
4		Eve	50	Other	48000	eve@email.com	Company_A
5		Frank	Unknown	Male	N/A		Company_A
6		Grace	30	Female	51000	grace@email.com	Company_A
7		Alice	25	Female	50000	alice@email.com	Company_A

	Name	Age	Gender	Salary	Email	DataGroup
0	Graham	85	Male	12000	g_able@email.com	Company_B
1	Simon	76	Male	52000	simon@here.com	Company_B
2	John	25	NaN	55000	jj@email.com	Company_B
3	Fred	43	Male	48000	fred_clampett@email.com	Company_B
4	Francis	62	Female	50000	frossi@quo.com	Company_B

```
# df_concat = pd.concat([df1,df2], ignore_index = True, sort=True)
print(df_concat)
```

	Age	DataGroup	Email	Gender	Name	Salary	
0	25	Company_A	alice@email.com	Female	Alice	50000	
1	NaN	Company_A	bob@email	Male	Bob	Newhart	52000
2	35	Company_A	charlie--email.com	NaN	Charlie	55000	
3	40	Company_A		Female	NaN	NaN	NaN
4	50	Company_A	eve@email.com	Other	Eve	48000	
5	Unknown	Company_A		Male	Frank	N/A	
6	30	Company_A	grace@email.com	Female	Grace	51000	
7	25	Company_A	alice@email.com	Female	Alice	50000	
8	85	Company_B	g_able@email.com	Male	Graham	12000	
9	76	Company_B	simon@here.com	Male	Simon	52000	
10	25	Company_B	jj@email.com	NaN	John	55000	
11	43	Company_B	fred_clampett@email.com	Male	Fred	48000	
12	62	Company_B	frossi@quo.com	Female	Francis	50000	

EXERCISE 11:

USE THE CONCATENATE TO ADD A NEW ROW



- Add a newdata row using a list converted to a dataframe as:

```
newdata = {  
    'station_id': ['Downtown'], 'name': ['abc'], 'lat': ['47.7'], 'long': ['-122.3']}  
df_cycle_stationNew = pd.DataFrame(newdata)
```


Use concat to add the newdata to the current dataframe

- Review the contents of the dataframe

write to csv file

Write to csv

```
df_stats.to_csv("stats.csv",sep=',')
```



EXERCISE 12: WRITE TO A CSV



Using the describe function from the last exercise output a new database (df_stats)


Write the dt_stats dataframe to a csv file

Matplotlib

Add a chart to the notebook using matplotlib

```
In [1]: from matplotlib import pyplot as plt  
plt.bar(x=df['Age'], height=df['Salary'])
```

```
[1]: <BarContainer object of 7 artists>
```



Updating styles

```
# Clear the plot area  
plt.clf()  
  
# Create a bar plot of revenue by year  
plt.bar(x=df['Age'], height=df['Salary'], color='orange')  
  
# Customize the chart  
plt.title('Salary by year')  
plt.xlabel('Year')  
plt.ylabel('Salary')  
plt.grid(color='#95a5a6', linestyle='--', linewidth=2, axis='y', alpha=0.7)  
plt.xticks(rotation=45)  
  
# Show the figure  
plt.show()
```




EXERCISE 13:

CREATE A PLOT USING MATPLOTLIB




- Write the code to display a chart similar to the picture below
- Review the matplotlib help pages on the xticks to show 0 to 32 inclusive



EXAMPLE CODE


```
▶ import pandas as pd
    import matplotlib.pyplot as plt

    grouped = df_cycle_station.groupby(['install_dockcount']).size().reset_index(name='count')
    grouped.dtypes
    plt.bar(x=grouped['install_dockcount'],height=grouped['count'])
    plt.xticks([0,2,4,6,8,10,12,14,16,18,20,22,24,26,28,30,32])
    plt.show()
```



Using seaborn

```
# import seaborn as sns  
  
# Clear the plot area  
plt.clf()  
# Create a bar chart  
ax = sns.barplot(x="Age", y="Salary", data=df)  
plt.show()
```



EXERCISE 14:

USING SEABORN



- Clear the current plot
- Create a barplot within seaborn